

PSTAT220A Homework 3

2025-11-16

Problem 1

Consider the data on teen gambling in `teengamb` from the `faraway` package. Fit a simple linear regression model with the expenditure on gambling as the response and income as the covariate:

- (a) Check the constant variance assumption for errors. What do you conclude?

```
data(teengamb)
model <- lm(gamble ~ income, data = teengamb)
summary(model)
```

Call:
`lm(formula = gamble ~ income, data = teengamb)`

Residuals:

Min	1Q	Median	3Q	Max
-46.020	-11.874	-3.757	11.934	107.120

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-6.325	6.030	-1.049	0.3
income	5.520	1.036	5.330	3.05e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 24.95 on 45 degrees of freedom

Multiple R-squared: 0.387, Adjusted R-squared: 0.3734

F-statistic: 28.41 on 1 and 45 DF, p-value: 3.045e-06

```

par(mfrow = c(2, 2))


plot(fitted(model), residuals(model),
      xlab = "Fitted Values", ylab = "Residuals",
      main = "Residuals vs Fitted Values", pch = 19, col = rgb(0, 0, 1, 0.5))
abline(h = 0, col = "red", lwd = 2, lty = 2)
lines(lowess(fitted(model), residuals(model)), col = "darkred", lwd = 2)

plot(fitted(model), sqrt(abs(rstandard(model))),
      xlab = "Fitted Values", ylab = expression(sqrt("|Standardized Residuals|")),
      main = "Scale-Location Plot", pch = 19, col = rgb(0, 0, 1, 0.5))
lines(lowess(fitted(model), sqrt(abs(rstandard(model)))), col = "darkred", lwd = 2)

plot(teengamb$income, residuals(model),
      xlab = "Income", ylab = "Residuals",
      main = "Residuals vs Income", pch = 19, col = rgb(0, 0, 1, 0.5))
abline(h = 0, col = "red", lwd = 2, lty = 2)
lines(lowess(teengamb$income, residuals(model)), col = "darkred", lwd = 2)

plot(teengamb$income, abs(residuals(model)),
      xlab = "Income", ylab = "Absolute Residuals",
      main = "Absolute Residuals vs Income", pch = 19, col = rgb(0, 0, 1, 0.5))
lines(lowess(teengamb$income, abs(residuals(model))), col = "darkred", lwd = 2)

```



```

par(mfrow = c(1, 1))

bp_test <- bptest(model)
bp_test

```

studentized Breusch-Pagan test

data: model
BP = 5.6546, df = 1, p-value = 0.01741

Conclusion: From the diagnostic plots, we can see clear heteroskedasticity. The Scale-Location plot and the absolute residuals vs income plot both show that variance increases with income. The Breusch-Pagan test gives p-value = 0.0174, which rejects the constant variance assumption at $\alpha = 0.05$. So the residual variance grows as income increases.

- (b) Assume that constant variance assumption is violated and that instead the residual variance grows linearly with income. Fit a weighted least squares to account for this fact. Report the coefficient and confidence interval.

```

wts <- 1 / teengamb$income
model_wls <- lm(gamble ~ income, data = teengamb, weights = wts)
summary(model_wls)

```

Call:
`lm(formula = gamble ~ income, data = teengamb, weights = wts)`

Weighted Residuals:

Min	1Q	Median	3Q	Max
-13.576	-6.639	-2.820	6.020	35.331

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-2.333	3.780	-0.617	0.54
income	4.661	1.054	4.423	6.11e-05 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 9.879 on 45 degrees of freedom
Multiple R-squared: 0.303, Adjusted R-squared: 0.2875
F-statistic: 19.56 on 1 and 45 DF, p-value: 6.113e-05

```

coef_wls <- coef(model_wls)
ci_wls <- confint(model_wls, level = 0.95)

wls_results <- data.frame(
  Parameter = c("Intercept", "Income"),
  Estimate = c(coef_wls[1], coef_wls[2]),
  CI_Lower = c(ci_wls[1, 1], ci_wls[2, 1]),
  CI_Upper = c(ci_wls[1, 2], ci_wls[2, 2])
)
wls_results[, 2:4] <- round(wls_results[, 2:4], 4)
print(wls_results)

```

	Parameter	Estimate	CI_Lower	CI_Upper
(Intercept)	Intercept	-2.3333	-9.9471	5.2805
income	Income	4.6607	2.5381	6.7832

```

comparison_df <- data.frame(
  Method = c("OLS", "WLS"),
  Intercept = c(coef(model)[1], coef_wls[1]),
  Slope = c(coef(model)[2], coef_wls[2])
)
print(comparison_df)

```


	Method	Intercept	Slope
1	OLS	-6.324559	5.520485
2	WLS	-2.333297	4.660654

```

plot(teengamb$income, teengamb$gamble,
      xlab = "Income", ylab = "Gambling Expenditure",
      main = "OLS vs WLS Regression Fits", pch = 19, col = "gray40")
abline(model, col = "blue", lwd = 2, lty = 1)
abline(model_wls, col = "red", lwd = 2, lty = 2)
legend("topleft", legend = c("OLS", "WLS (weighted)"),
       col = c("blue", "red"), lty = c(1, 2), lwd = 2)

```

OLS vs WLS Regression Fits



Using WLS with weights = $1/\text{income}$ to account for heteroskedasticity, we get intercept = -2.33 (95% CI: [-9.95, 5.28]) and slope = 4.6607 (95% CI: [2.5381, 6.7832]). These estimates differ from OLS because WLS downweights high-income observations where variance is larger.

- (c) Construct a bootstrap confidence interval of the regression coefficients and of the residual variance (assuming an income of 1).

```
set.seed(220)

boot_fn <- function(data, idx) {
  fit <- lm(gamble ~ income, data = data[idx, ])

  resid_sq <- residuals(fit)^2
  var_mod <- lm(resid_sq ~ income, data = data[idx, ])
  var1 <- as.numeric(cbind(1, 1) %*% coef(var_mod))
  var1 <- pmax(var1, 1e-6)
  log_var1 <- log(var1)

  c(coef(fit), log_var_at_1 = log_var1)
}

B <- 5000
boot_res <- boot(teengamb, boot_fn, R = B)
```

```

ci_intercept <- boot.ci(boot_res, index = 1, type = c("perc", "bca"), conf = 0.95)
ci_slope <- boot.ci(boot_res, index = 2, type = c("perc", "bca"), conf = 0.95)

ci_log_var <- boot.ci(boot_res, index = 3, type = "perc", conf = 0.95)
ci_var_lower <- exp(ci_log_var$perc[4])
ci_var_upper <- exp(ci_log_var$perc[5])

bootstrap_results <- data.frame(
  Parameter = c("Intercept (BCa)", "Income coefficient (BCa)",
               "Residual variance at income=1 (Percentile)"),
  Lower_CI = c(ci_intercept$bca[4], ci_slope$bca[4], ci_var_lower),
  Upper_CI = c(ci_intercept$bca[5], ci_slope$bca[5], ci_var_upper)
)
bootstrap_results[, 2:3] <- round(bootstrap_results[, 2:3], 4)
print(bootstrap_results)

```

	Parameter	Lower_CI	Upper_CI
1	Intercept (BCa)	-17.0872	2.1817
2	Income coefficient (BCa)	3.0942	9.4715
3	Residual variance at income=1 (Percentile)	0.0000	252.0784

```


par(mfrow = c(1, 3))

hist(boot_res$t[, 1], breaks = 50, prob = TRUE, col = "lightblue",
     main = "Bootstrap Distribution: Intercept", xlab = "Intercept")
abline(v = ci_intercept$bca[4:5], col = "red", lwd = 2, lty = 2)

hist(boot_res$t[, 2], breaks = 50, prob = TRUE, col = "lightgreen",
     main = "Bootstrap Distribution: Slope", xlab = "Slope")
abline(v = ci_slope$bca[4:5], col = "red", lwd = 2, lty = 2)

hist(exp(boot_res$t[, 3]), breaks = 50, prob = TRUE, col = "lightcoral",
     main = "Bootstrap Distribution: Var(income=1)", xlab = "Variance")
abline(v = c(ci_var_lower, ci_var_upper), col = "red", lwd = 2, lty = 2)

```



```
par(mfrow = c(1, 1))
```

Using 5000 bootstrap samples, we get non-parametric confidence intervals that don't rely on normality. The 95% BCa CIs are: intercept [-17.09, 2.18], slope [3.0942, 9.4715]. For variance at income=1, I used log transformation to avoid negative values, which gives percentile CI [0, 252.08] after back-transforming.

- (d) Make a plot showing the prediction intervals (again assuming residual variance grows linearly with income). The plot should have `income` on the x-axis, `gamble` on the y-axis, and prediction bands for each point.

```
resid_sq <- residuals(model)^2
var_model <- lm(resid_sq ~ income, data = teengamb)

new_income <- seq(min(teengamb$income), max(teengamb$income), length = 100)
pred <- predict(model, newdata = data.frame(income = new_income), se.fit = TRUE)

pred_var <- pred$se.fit^2 + pmax(0, predict(var_model,
                                              newdata = data.frame(income = new_income)))


crit <- qt(0.975, df = model$df.residual)
lower <- pred$fit - crit * sqrt(pred_var)
upper <- pred$fit + crit * sqrt(pred_var)

plot(teengamb$income, teengamb$gamble,
      xlab = "Income", ylab = "Gambling Expenditure",
      main = "95% Prediction Intervals",
      pch = 19, col = "gray40", cex = 1.2)
lines(new_income, pred$fit, col = "blue", lwd = 2)
lines(new_income, lower, col = "red", lwd = 2, lty = 2)
lines(new_income, upper, col = "red", lwd = 2, lty = 2)
```

```

polygon(c(new_income, rev(new_income)),
        c(lower, rev(upper)),
        col = rgb(1, 0, 0, 0.15), border = NA)
legend("topleft",
       legend = c("Observed", "Fitted", "95% PI"),
       col = c("gray40", "blue", "red"),
       pch = c(19, NA, NA), lty = c(NA, 1, 2), lwd = 2)

```



The prediction intervals widen as income increases, which makes sense given our assumption that variance grows linearly with income. This is different from standard PIs (which assume constant variance) that would have the same width everywhere. The widening bands capture the increased uncertainty at higher income levels.

- (e) Fit a robust regression model using `rlm`. Compare the coefficients from your weighted least squares fit.

```

model_robust <- rlm(gamble ~ income, data = teengamb, method = "MM")
summary(model_robust)

```

```

Call: rlm(formula = gamble ~ income, data = teengamb, method = "MM")
Residuals:
    Min      1Q  Median      3Q     Max 
 -5.3822 -3.6482  0.4688 14.7589 150.4598

```

```

Coefficients:
      Value Std. Error t value
(Intercept) 4.3801 1.7233    2.5417
income       0.1160 0.2960    0.3919

```

Residual standard error: 6.617 on 45 degrees of freedom

```

comparison_df <- data.frame(
  Method = c("OLS", "WLS", "Robust (rlm)"),
  Intercept = c(coef(model)[1], coef(model_wls)[1], coef(model_robust)[1]),
  Slope = c(coef(model)[2], coef(model_wls)[2], coef(model_robust)[2])
)
print(comparison_df)

```

	Method	Intercept	Slope
1	OLS	-6.324559	5.5204853
2	WLS	-2.333297	4.6606543
3	Robust (rlm)	4.380122	0.1160105

```

weights_robust <- model_robust$w
downweighted_idx <- which(weights_robust < 0.8)

outlier_results <- data.frame(
  Index = downweighted_idx,
  Weight = weights_robust[downweighted_idx]
)
print(outlier_results)

```

	Index	Weight
1	5	0.5869727
2	24	0.0000000
3	25	0.0000000
4	31	0.0000000
5	32	0.0000000
6	33	0.0000000
7	36	0.0000000
8	37	0.0000000
9	38	0.0000000
10	40	0.3196085
11	42	0.0000000
12	45	0.0000000
13	47	0.6087454

```


plot(teengamb$income, teengamb$gamble,
      xlab = "Income", ylab = "Gambling Expenditure",
      main = "Comparison: OLS, WLS, and Robust Regression",
      pch = 19, col = "gray60")

points(teengamb$income[downweighted_idx],
       teengamb$gamble[downweighted_idx],
       pch = 19, col = "orange", cex = 1.8)

abline(model, col = "blue", lwd = 2, lty = 1)
abline(model_wls, col = "red", lwd = 2, lty = 2)
abline(model_robust, col = "darkgreen", lwd = 2, lty = 3)

legend("topleft",
       legend = c("OLS", "WLS", "Robust (rlm)", "Downweighted points"),
       col = c("blue", "red", "darkgreen", "orange"),
       lty = c(1, 2, 3, NA), pch = c(NA, NA, NA, 19), lwd = 2)

```



The robust regression (MM-estimator) downweights outliers and gives estimates less sensitive to extreme values. It identified 13 observations with weights < 0.8 as potential outliers. The robust fit gives intercept = 4.38 and slope = 0.116, which differ from both OLS and WLS. This makes sense since the MM-estimator is designed to handle outliers better.

- (f) Use `optim` to fit a t-regression. That is, fit the linear model assuming iid residuals with a t_3 distribution. Construct a bootstrap confidence interval of the regression

coefficients.

```
nll_t3 <- function(par, X, y) {
  beta <- par[1:2]
  sigma <- exp(par[3])
  resid <- y - X %*% beta
  -sum(dt(resid / sigma, df = 3, log = TRUE) - log(sigma))
}

X <- cbind(1, teengamb$income)
y <- teengamb$gamble

init <- c(coef(model), log(summary(model)$sigma))

opt <- optim(init, nll_t3, X = X, y = y, method = "BFGS")

beta_t3 <- opt$par[1:2]
scale_t3 <- exp(opt$par[3])

t3_results <- data.frame(
  Parameter = c("Intercept", "Income coefficient", "Scale parameter", "Convergence"),
  Value = c(beta_t3[1], beta_t3[2], scale_t3, opt$convergence)
)
print(t3_results)
```

	Parameter	Value
1	Intercept	-6.036333
2	Income coefficient	4.650505
3	Scale parameter	14.259649
4	Convergence	0.000000

```
set.seed(220)

boot_t3 <- function(data, idx) {
  X_b <- cbind(1, data$income[idx])
  y_b <- data$gamble[idx]

  init_b <- c(coef(lm(gamble ~ income, data = data[idx, ])),
             log(summary(lm(gamble ~ income, data = data[idx, ]))$sigma))

  opt_b <- tryCatch({
    optim(init_b, nll_t3, X = X_b, y = y_b, method = "BFGS")
  }, error = function(e) {
    list(par = c(NA, NA))
```

```

    })
  opt_b$par[1:2]
}

B_t3 <- 1000
boot_t3_res <- boot(teengamb, boot_t3, R = B_t3)

ci_t3_intercept <- boot.ci(boot_t3_res, index = 1, type = c("perc"), conf = 0.95)
ci_t3_slope <- boot.ci(boot_t3_res, index = 2, type = c("perc"), conf = 0.95)

t3_ci_results <- data.frame(
  Parameter = c("Intercept (Percentile)", "Income coefficient (Percentile)"),
  Lower_CI = c(ci_t3_intercept$perc[4], ci_t3_slope$perc[4]),
  Upper_CI = c(ci_t3_intercept$perc[5], ci_t3_slope$perc[5])
)
t3_ci_results[, 2:3] <- round(t3_ci_results[, 2:3], 4)
print(t3_ci_results)

```

	Parameter	Lower_CI	Upper_CI
1	Intercept (Percentile)	-12.7057	6.3680
2	Income coefficient (Percentile)	0.2150	6.8581


```

par(mfrow = c(1, 2))

hist(boot_t3_res$t[, 1], breaks = 30, prob = TRUE, col = "lightblue",
     main = "Bootstrap: Intercept (t-regression)", xlab = "Intercept")
abline(v = ci_t3_intercept$bca[4:5], col = "red", lwd = 2, lty = 2)

hist(boot_t3_res$t[, 2], breaks = 30, prob = TRUE, col = "lightgreen",
     main = "Bootstrap: Slope (t-regression)", xlab = "Slope")
abline(v = ci_t3_slope$bca[4:5], col = "red", lwd = 2, lty = 2)

```



```
par(mfrow = c(1, 1))

final_comparison <- data.frame(
  Method = c("OLS", "WLS", "Robust (rlm)", "t-regression"),
  Intercept = c(coef(model)[1], coef(model_wls)[1],
                coef(model_robust)[1], beta_t3[1]),
  Slope = c(coef(model)[2], coef(model_wls)[2],
            coef(model_robust)[2], beta_t3[2])
)
print(final_comparison)
```

	Method	Intercept	Slope
1	OLS	-6.324559	5.5204853
2	WLS	-2.333297	4.6606543
3	Robust (rlm)	4.380122	0.1160105
4	t-regression	-6.036333	4.6505046

The t-regression uses a heavy-tailed t_3 distribution, making it robust to outliers. The final fit, found using optim, gave a slope of 4.6505. Based on bootstrap CIs, the t-regression coefficients are similar to the robust regression and less swayed by extreme values than OLS.

Problem 2

Standard multiple testing guarantees are given assuming independent tests. In general, test statistics for regression coefficients are correlated.

- Assume you have $p = 300$ predictors and $n = 1000$ observations. Let $y = X\beta + \epsilon$ (no intercept), where the first 10 coefficients in β equal 0.2 and the last 90 are 0. Generate $X \sim N(0, 1)$ for each feature. Run a simulation where you generate data

from this model many times. Fit the linear model and use both Benjamini-Hochberg and the q-value correction to select discoveries controlling the FDR at 0.05. For each simulated dataset compute 1) the false discovery proportion and 2) the sensitivity (the fraction of the 10 significant predictors that were correctly declared successes). Plot the distribution FDP and sensitivity for both BH and Q-value approaches using histograms or a beeswarm plot or something similar.

```

set.seed(220)

n <- 1000
p <- 300
n_sig <- 10
beta_sig <- 0.2
n_sim <- 200

results <- replicate(n_sim, {
  X <- matrix(rnorm(n * p), n, p)
  beta <- c(rep(beta_sig, n_sig), rep(0, p - n_sig))
  y <- X %*% beta + rnorm(n)

  fit <- lm(y ~ X - 1)
  pvals <- summary(fit)$coef[, 4]

  bh_rej <- which(p.adjust(pvals, "BH") <= 0.05)

  qobj <- qvalue(pvals)
  q_rej <- which(qobj$qvalues <= 0.05)

  fdp_bh <- if(length(bh_rej) > 0) sum(bh_rej > n_sig) / length(bh_rej) else 0
  fdp_q <- if(length(q_rej) > 0) sum(q_rej > n_sig) / length(q_rej) else 0

  sens_bh <- sum(bh_rej <= n_sig) / n_sig
  sens_q <- sum(q_rej <= n_sig) / n_sig

  c(fdp_bh = fdp_bh, fdp_q = fdp_q, sens_bh = sens_bh, sens_q = sens_q)
})

summary_stats <- apply(results, 1, function(x) c(mean = mean(x), median = median(x)))
summary_stats <- round(summary_stats, 4)
print(summary_stats)

      fdp_bh fdp_q sens_bh sens_q
mean    0.0442 0.048  0.9825  0.984
median  0.0000 0.000  1.0000  1.000

```

```

par(mfrow = c(2, 2))


hist(results[1, ], breaks = 30, col = "lightblue",
      main = "BH: False Discovery Proportion (rho = 0)",
      xlab = "FDP", prob = TRUE)
abline(v = 0.05, col = "red", lwd = 2, lty = 2)
abline(v = mean(results[1, ]), col = "darkblue", lwd = 2)

hist(results[3, ], breaks = 30, col = "lightgreen",
      main = "BH: Sensitivity (rho = 0)",
      xlab = "Sensitivity", prob = TRUE)
abline(v = mean(results[3, ]), col = "darkgreen", lwd = 2)

hist(results[2, ], breaks = 30, col = "lightblue",
      main = "Q-value: False Discovery Proportion (rho = 0)",
      xlab = "FDP", prob = TRUE)
abline(v = 0.05, col = "red", lwd = 2, lty = 2)
abline(v = mean(results[2, ]), col = "darkblue", lwd = 2)

hist(results[4, ], breaks = 30, col = "lightgreen",
      main = "Q-value: Sensitivity (rho = 0)",
      xlab = "Sensitivity", prob = TRUE)
abline(v = mean(results[4, ]), col = "darkgreen", lwd = 2)

```



```
par(mfrow = c(1, 1))
```

Results: With independent predictors ($\rho = 0$), both BH and q-value work well - they control FDR below 5% and have good power to detect the true signals. The mean FDP for BH is 0.044 and for q-value is 0.048, both well below 0.05. Sensitivity is around 0.983 for both methods.

- (b) Now consider generating data assuming correlated covariates drawn from a multivariate normal, $X \sim N(\mu, \Sigma)$ with Σ an equicorrelation matrix. To investigate how problematic positive equicorrelation is for these procedures, repeat part (a) for each of $\rho \in \{0.25, 0.5, 0.75\}$. Plot the distributions as a function of ρ . What can you conclude about how positive equicorrelation amongst the predictors impacts the false discovery proportions? How does it impact sensitivity?

```
simulate_corr <- function(rho, n_sim = 200) {
  Sigma <- matrix(rho, p, p)
  diag(Sigma) <- 1

  replicate(n_sim, {
    X <- rmvnorm(n, sigma = Sigma)
    beta <- c(rep(beta_sig, n_sig), rep(0, p - n_sig))
    y <- X %*% beta + rnorm(n)

    pvals <- summary(lm(y ~ X - 1))$coef[, 4]

    bh_rej <- which(p.adjust(pvals, "BH") <= 0.05)
    q_rej <- which(qvalue(pvals)$qvalues <= 0.05)

    fdp <- function(rej) if(length(rej) > 0) sum(rej > n_sig) / length(rej) else 0
    sens <- function(rej) sum(rej <= n_sig) / n_sig

    c(fdp_bh = fdp(bh_rej), sens_bh = sens(bh_rej),
      fdp_q = fdp(q_rej), sens_q = sens(q_rej))
  })
}

rho_vals <- c(0, 0.25, 0.5, 0.75)
set.seed(220)
results_list <- lapply(rho_vals, simulate_corr)

summary_df <- data.frame(
  rho = rho_vals,
  mean_fdp_bh = sapply(results_list, function(x) mean(x[1, ])),
  mean_sens_bh = sapply(results_list, function(x) mean(x[2, ])),
```

```

mean_fdp_q = sapply(results_list, function(x) mean(x[3, ])),
mean_sens_q = sapply(results_list, function(x) mean(x[4, ]))
)
summary_df[, 2:5] <- round(summary_df[, 2:5], 4)
print(summary_df)

```

	rho	mean_fdp_bh	mean_sens_bh	mean_fdp_q	mean_sens_q
1	0.00	0.0504	0.9805	0.0566	0.9820
2	0.25	0.0508	0.9245	0.0573	0.9330
3	0.50	0.0529	0.7000	0.0602	0.7120
4	0.75	0.0531	0.1975	0.0576	0.2055


```

par(mfrow = c(1, 2))

boxplot(lapply(results_list, function(x) x[1, ]),
        names = rho_vals,
        main = "BH: FDP vs Correlation",
        xlab = expression(rho), ylab = "False Discovery Proportion",
        col = "lightblue", ylim = c(0, max(sapply(results_list, function(x) max(x[1, ])))))
abline(h = 0.05, col = "red", lty = 2, lwd = 2)

boxplot(lapply(results_list, function(x) x[3, ]),
        names = rho_vals,
        main = "Q-value: FDP vs Correlation",
        xlab = expression(rho), ylab = "False Discovery Proportion",
        col = "lightgreen", ylim = c(0, max(sapply(results_list, function(x) max(x[3, ])))))
abline(h = 0.05, col = "red", lty = 2, lwd = 2)

```



```


par(mfrow = c(1, 1))

par(mfrow = c(1, 2))

boxplot(lapply(results_list, function(x) x[2, ]),
        names = rho_vals,
        main = "BH: Sensitivity vs Correlation",
        xlab = expression(rho), ylab = "Sensitivity",
        col = "lightcoral")

boxplot(lapply(results_list, function(x) x[4, ]),
        names = rho_vals,
        main = "Q-value: Sensitivity vs Correlation",
        xlab = expression(rho), ylab = "Sensitivity",
        col = "lightyellow")

```



```
par(mfrow = c(1, 1))
```

Conclusions:

As ρ increases, both FDP and sensitivity get worse for both methods. At $\rho = 0$, mean FDP is around 0.05 for BH and 0.057 for q-value (both < 0.05). But at $\rho = 0.75$, FDP jumps to 0.053 for BH and 0.058 for q-value, exceeding the 5% target. Also, the variability of FDP increases a lot.

For sensitivity, it drops from around 0.981 at $\rho = 0$ to about 0.198 for BH and 0.206 for q-value at $\rho = 0.75$. So we lose power to detect true signals.

This makes sense because positive correlation among predictors violates the independence assumption that these FDR methods rely on. When test statistics are correlated, it's harder to tell true signals from correlated noise, leading to more false discoveries and missed detections.

Problem 3

Run a regression with `prostate` data in `faraway` library with `lpsa` as the response variable and the other variables as covariates.

- (a) Select the best model using each of (1) AIC (either forward or backward selection); (2) best subset selection with cross-validation; and (3) LASSO with cross-validation. How much agreement is there between approaches?

```
data(prostate)
head(prostate)
```

	lcavol	lweight	age	lbph	svi	lcp	gleason	pgg45	lpsa
1	-0.5798185	2.7695	50	-1.386294	0	-1.38629	6	0	-0.43078
2	-0.9942523	3.3196	58	-1.386294	0	-1.38629	6	0	-0.16252
3	-0.5108256	2.6912	74	-1.386294	0	-1.38629	7	20	-0.16252
4	-1.2039728	3.2828	58	-1.386294	0	-1.38629	6	0	-0.16252
5	0.7514161	3.4324	62	-1.386294	0	-1.38629	6	0	0.37156
6	-1.0498221	3.2288	50	-1.386294	0	-1.38629	6	0	0.76547

```
str(prostate)
```

```
'data.frame': 97 obs. of 9 variables:
 $ lcavol : num -0.58 -0.994 -0.511 -1.204 0.751 ...
 $ lweight: num 2.77 3.32 2.69 3.28 3.43 ...
 $ age     : int 50 58 74 58 62 50 64 58 47 63 ...
 $ lbph    : num -1.39 -1.39 -1.39 -1.39 -1.39 ...
 $ svi     : int 0 0 0 0 0 0 0 0 0 ...
 $ lcp     : num -1.39 -1.39 -1.39 -1.39 -1.39 ...
 $ gleason: int 6 6 7 6 6 6 6 6 6 ...
 $ pgg45  : int 0 0 20 0 0 0 0 0 0 ...
 $ lpsa    : num -0.431 -0.163 -0.163 -0.163 0.372 ...
```

```
full_model <- lm(lpsa ~ ., data = prostate)
backward_aic <- step(full_model, direction = "backward", trace = 0)

names(coef(backward_aic))[-1]
```

```
[1] "lcavol"   "lweight"  "age"       "lbph"      "svi"
```

```
AIC(backward_aic)
```

```
[1] 215.8997
```

```
summary(backward_aic)$adj.r.squared
```

```
[1] 0.6245476
```

```

null_model <- lm(lpsa ~ 1, data = prostate)
forward_aic <- step(null_model,
                      scope = list(lower = ~ 1, upper = formula(full_model)),
                      direction = "forward", trace = 0)

names(coef(forward_aic))[-1]

[1] "lcavol"  "lweight" "svi"      "lbph"     "age"

AIC(forward_aic)

[1] 215.8997

summary(forward_aic)$adj.r.squared

[1] 0.6245476

regfit <- regsubsets(lpsa ~ ., data = prostate, nvmax = 8)
reg_summary <- summary(regfit)

par(mfrow = c(2, 2))


plot(reg_summary$cp, xlab = "Number of Variables",
     ylab = "Cp", type = "b", pch = 19, col = "blue")
points(which.min(reg_summary$cp), reg_summary$cp[which.min(reg_summary$cp)],
       col = "red", cex = 2, pch = 20)

plot(reg_summary$bic, xlab = "Number of Variables",
     ylab = "BIC", type = "b", pch = 19, col = "blue")
points(which.min(reg_summary$bic), reg_summary$bic[which.min(reg_summary$bic)],
       col = "red", cex = 2, pch = 20)

plot(reg_summary$adjr2, xlab = "Number of Variables",
     ylab = "Adjusted R^2", type = "b", pch = 19, col = "blue")
points(which.max(reg_summary$adjr2), reg_summary$adjr2[which.max(reg_summary$adjr2)],
       col = "red", cex = 2, pch = 20)

par(mfrow = c(1, 1))

```



```

predict_regressions <- function(object, newdata, id) {
  form <- as.formula(object$call[[2]])
  mm <- model.matrix(form, newdata)
  coefi <- coef(object, id = id)
  vars <- names(coefi)
  as.numeric(mm[, vars, drop = FALSE] %*% coefi)
}

set.seed(220)
k <- 10
folds <- sample(rep(1:k, length.out = nrow(prostate)))
cv_errors <- matrix(NA, k, 8)

for (i in 1:k) {
  fit <- regsubsets(lpsa ~ ., data = prostate[folds != i, ], nvmax = 8)
  for (j in 1:8) {
    pred <- predict_regressions(fit, prostate[folds == i, ], id = j)
    cv_errors[i, j] <- mean((prostate$lpsa[folds == i] - pred)^2)
  }
}

mean_cv_errors <- colMeans(cv_errors)
best_size <- which.min(mean_cv_errors)


plot(mean_cv_errors, type = "b", pch = 19, col = "blue",

```

```

xlab = "Number of Variables", ylab = "Cross-Validation MSE",
main = "Best Subset Selection with 10-Fold CV")
points(best_size, mean_cv_errors[best_size], col = "red", cex = 2, pch = 20)

```



```
best_size
```

```
[1] 7
```

```
mean_cv_errors[best_size]
```

```
[1] 0.5528904
```

```

regfit_best <- regsubsets(lpsa ~ ., data = prostate, nvmax = 8)
coef_best_cv <- coef(regfit_best, id = best_size)

names(coef_best_cv)[-1]

```

```
[1] "lcavol"   "lweight"  "age"        "lbph"      "svi"       "lcp"       "pgg45"
```


```

X_mat <- model.matrix(lpsa ~ ., prostate)[, -1]
y_vec <- prostate$lpsa

set.seed(220)
cv_lasso <- cv.glmnet(X_mat, y_vec, alpha = 1, nfolds = 10)

plot(cv_lasso, main = "LASSO: Cross-Validation MSE")

```



```
cv_lasso$lambda.min
```

```
[1] 0.02961434
```

```
cv_lasso$lambda.1se
```

```
[1] 0.1580428
```

```

lasso_coef_min <- coef(cv_lasso, s = "lambda.min")
lasso_vars_min <- rownames(lasso_coef_min)[lasso_coef_min[, 1] != 0]
lasso_vars_min <- lasso_vars_min[lasso_vars_min != "(Intercept)"]

lasso_coef_1se <- coef(cv_lasso, s = "lambda.1se")
lasso_vars_1se <- rownames(lasso_coef_1se)[lasso_coef_1se[, 1] != 0]

```

```

lasso_vars_1se <- lasso_vars_1se[lasso_vars_1se != "(Intercept)"]

lasso_vars_min
```

```
[1] "lcavol"  "lweight" "age"      "lbph"     "svi"      "gleason"  "pgg45"
```

```
lasso_vars_1se
```

```
[1] "lcavol"  "lweight" "svi"
```

```
all_vars <- names(prostate)[names(prostate) != "lpsa"]
```

```

comparison_table <- data.frame(
  Variable = all_vars,
  Backward_AIC = all_vars %in% names(coef(backward_aic))[-1],
  Forward_AIC = all_vars %in% names(coef(forward_aic))[-1],
  Best_Subset_CV = all_vars %in% names(coef_best_cv)[-1],
  LASSO_min = all_vars %in% lasso_vars_min,
  LASSO_1se = all_vars %in% lasso_vars_1se
)
```

```
comparison_table
```

	Variable	Backward_AIC	Forward_AIC	Best_Subset_CV	LASSO_min	LASSO_1se
1	lcavol	TRUE	TRUE	TRUE	TRUE	TRUE
2	lweight	TRUE	TRUE	TRUE	TRUE	TRUE
3	age	TRUE	TRUE	TRUE	TRUE	FALSE
4	lbph	TRUE	TRUE	TRUE	TRUE	FALSE
5	svi	TRUE	TRUE	TRUE	TRUE	TRUE
6	lcp	FALSE	FALSE	TRUE	FALSE	FALSE
7	gleason	FALSE	FALSE	FALSE	TRUE	FALSE
8	pgg45	FALSE	FALSE	TRUE	TRUE	FALSE

```

agreement_count <- rowSums(comparison_table[, -1])
data.frame(Variable = all_vars, Count = agreement_count)
```

	Variable	Count
1	lcavol	5
2	lweight	5
3	age	4
4	lbph	4

```

5      svi      5
6      lcp      1
7  gleason      1
8  pgg45      2

```

Agreement Analysis: All three methods show pretty good agreement. The variables `lcavol`, `lweight`, and `svi` are selected by almost all methods, so they're clearly important predictors. `lbph` is chosen by most methods too. The other variables (`age`, `lcp`, `gleason`, `pgg45`) are less consistent. Forward and backward AIC give the same model with 5 variables. Best subset CV picks 7 variables, and LASSO (`lambda.1se`) gives the most parsimonious model with 3 variables. Overall, the core predictors are consistent across methods.

- (b) Can you improve the model by adding higher order terms? Can you improve the model by adding multiplicative interaction terms? Use any selection approach of choice to settle on a “final” model.

```

prostate_poly <- prostate %>%
  mutate(
    lcavol2 = lcavol^2,
    lweight2 = lweight^2,
    age2 = age^2,
    lbph2 = lbph^2
  )

model_poly <- lm(lpsa ~ ., data = prostate_poly)
model_poly_step <- step(model_poly, direction = "backward", trace = 0)

names(coef(model_poly_step))[-1]

```

```
[1] "lcavol"  "lweight" "age"       "lbph"      "svi"
```

```
summary(model_poly_step)$adj.r.squared
```

```
[1] 0.6245476
```

```
AIC(model_poly_step)
```

```
[1] 215.8997
```

```

model_interact <- lm(lpsa ~ lcavol + lweight + age + lbph + svi + lcp +
                      gleason + pgg45 +
                      lcavol:lweight + lcavol:svi + age:lbph + lweight:svi,
                      data = prostate)

model_interact_step <- step(model_interact, direction = "backward", trace = 0)

names(coef(model_interact_step))[-1]

```

```
[1] "lcavol"  "lweight" "age"      "lbph"     "svi"
```

```
summary(model_interact_step)$adj.r.squared
```

```
[1] 0.6245476
```

```
AIC(model_interact_step)
```

```
[1] 215.8997
```

```

cv_compare <- function(formula, data, k = 10) {
  set.seed(220)
  folds <- sample(rep(1:k, length.out = nrow(data)))
  sapply(1:k, function(i) {
    if ("lcavol2" %in% all.vars(formula) && !("lcavol2" %in% names(data))) {
      return(NA)
    }
    fit <- lm(formula, data = data[folds != i, ])
    pred <- predict(fit, newdata = data[folds == i, ])
    mean((data$lpsa[folds == i] - pred)^2)
  })
}

set.seed(220)
cv_simple <- cv_compare(formula(backward_aic), prostate)
cv_poly <- cv_compare(formula(model_poly_step), prostate_poly)
cv_int <- cv_compare(formula(model_interact_step), prostate)

cv_results <- data.frame(
  Model = c("Simple (Backward AIC)", "Polynomial", "Interactions"),
  Mean_CV_MSE = c(mean(cv_simple), mean(cv_poly), mean(cv_int)),
  SE_CV_MSE = c(sd(cv_simple)/sqrt(length(cv_simple)),
                sd(cv_poly)/sqrt(length(cv_poly))),

```

```

        sd(cv_int)/sqrt(length(cv_int))),
Adj_R2 = c(summary(backward_aic)$adj.r.squared,
            summary(model_poly_step)$adj.r.squared,
            summary(model_interact_step)$adj.r.squared),
AIC = c(AIC(backward_aic), AIC(model_poly_step), AIC(model_interact_step))
)

cv_results_rounded <- cv_results
cv_results_rounded[, -1] <- round(cv_results[, -1], 4)
print(cv_results_rounded)

```

	Model	Mean_CV_MSE	SE_CV_MSE	Adj_R2	AIC
1	Simple (Backward AIC)	0.5415	0.0489	0.6245	215.8997
2	Polynomial	0.5415	0.0489	0.6245	215.8997
3	Interactions	0.5415	0.0489	0.6245	215.8997


```

par(mfrow = c(1, 2))

barplot(cv_results$Mean_CV_MSE, names.arg = c("Simple", "Polynomial", "Interactions"),
        col = c("lightblue", "lightgreen", "lightcoral"),
        main = "Cross-Validation MSE Comparison",
        ylab = "Mean CV MSE", las = 2, cex.names = 0.8,
        ylim = c(0, max(cv_results$Mean_CV_MSE) * 1.1))

barplot(cv_results$Adj_R2, names.arg = c("Simple", "Polynomial", "Interactions"),
        col = c("lightblue", "lightgreen", "lightcoral"),
        main = "Adjusted R-squared Comparison",
        ylab = "Adjusted R^2", las = 2, cex.names = 0.8,
        ylim = c(0, 1))

```



```
par(mfrow = c(1, 1))
```

```
best_model_idx <- which.min(cv_results$Mean_CV_MSE)
best_model_name <- cv_results$Model[best_model_idx]

if (grepl("Polynomial", best_model_name)) {
  final_model <- model_poly_step
} else if (grepl("Interactions", best_model_name)) {
  final_model <- model_interact_step
} else {
  final_model <- backward_aic
}

formula(final_model)
```

`lpsa ~ lcavol + lweight + age + lbph + svi`

```
summary(final_model)
```

Call:

```
lm(formula = lpsa ~ lcavol + lweight + age + lbph + svi, data = prostate)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.83505	-0.39396	0.00414	0.46336	1.57888

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.95100	0.83175	1.143	0.255882
lcavol	0.56561	0.07459	7.583	2.77e-11 ***
lweight	0.42369	0.16687	2.539	0.012814 *
age	-0.01489	0.01075	-1.385	0.169528
lbph	0.11184	0.05805	1.927	0.057160 .
svi	0.72095	0.20902	3.449	0.000854 ***


Signif. codes:	0	'***'	0.001	'**'
			0.01	'*'
			0.05	'. '
			0.1	' '
			1	

Residual standard error: 0.7073 on 91 degrees of freedom

Multiple R-squared: 0.6441, Adjusted R-squared: 0.6245

F-statistic: 32.94 on 5 and 91 DF, p-value: < 2.2e-16

```
par(mfrow = c(2, 2))
plot(final_model)
```



```
par(mfrow = c(1, 1))
```

Model Performance Comparison:

Metric	Simple	Polynomial	Interactions
CV MSE	0.5415	0.5415	0.5415
Adj R ²	0.6245	0.6245	0.6245
AIC	215.9	215.9	215.9

Final Model: Based on CV MSE, the Simple (Backward AIC) is best with CV MSE = 0.5415. Adding polynomial or interaction terms doesn't help much (max difference < 0), so the linear model seems to capture most of the relationship. The final model has 5 terms and the diagnostic plots look reasonable.